

Analytics 512 Homework 3

Arif Ali

02/21/16

Exercise 3.7 #10

```
In [1]: library(ISLR)
        attach(Carseats)
```

Part A

```
In [2]: Carseats.mm = lm(Sales ~ Price + Urban + US, data = Carseats)
```

Part B

```
In [3]: summary(Carseats.mm)
```

```
Out[3]: Call:
lm(formula = Sales ~ Price + Urban + US, data = Carseats)

Residuals:
    Min       1Q   Median       3Q      Max
-6.9206 -1.6220 -0.0564  1.5786  7.0581

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 13.043469   0.651012  20.036 < 2e-16 ***
Price       -0.054459   0.005242 -10.389 < 2e-16 ***
UrbanYes    -0.021916   0.271650  -0.081  0.936
USYes       1.200573    0.259042   4.635 4.86e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.472 on 396 degrees of freedom
Multiple R-squared:  0.2393,    Adjusted R-squared:  0.2335
F-statistic: 41.52 on 3 and 396 DF,  p-value: < 2.2e-16
```

Price: Price has a significant effect on the Response as evidenced by its P-value. However, the coefficient estimate is negative, which indicates a negative relationship.

UrbanYes: No relation is indicated based on the very high p-value. Based on a summary, UrbanYes is a qualitative variable indicating if a store is urban based or not.

USYes: USYes has a significant effect on the Response as evidenced by its P-value. The coefficient estimate indicates a positive relationship. USYes is a qualitative variable indicating if a store is based on the US.

Part C

$$\text{Sales} = 13.043469 - 0.054459 * \text{Price} - 0.021916 * \text{UrbanYes} + 1.200573 * \text{USYes}$$

Part D

Price and USYes based on the p-value.

Part E

```
In [4]: Carseats.mm.bs = lm(Sales ~ Price + US, data = Carseats)
summary(Carseats.mm.bs)
```

```
Out[4]: Call:
lm(formula = Sales ~ Price + US, data = Carseats)

Residuals:
    Min       1Q   Median       3Q      Max
-6.9269 -1.6286 -0.0574  1.5766  7.0515

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  13.03079    0.63098   20.652 < 2e-16 ***
Price       -0.05448    0.00523  -10.416 < 2e-16 ***
USYes        1.19964    0.25846    4.641 4.71e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.469 on 397 degrees of freedom
Multiple R-squared:  0.2393,    Adjusted R-squared:  0.2354
F-statistic: 62.43 on 2 and 397 DF,  p-value: < 2.2e-16
```

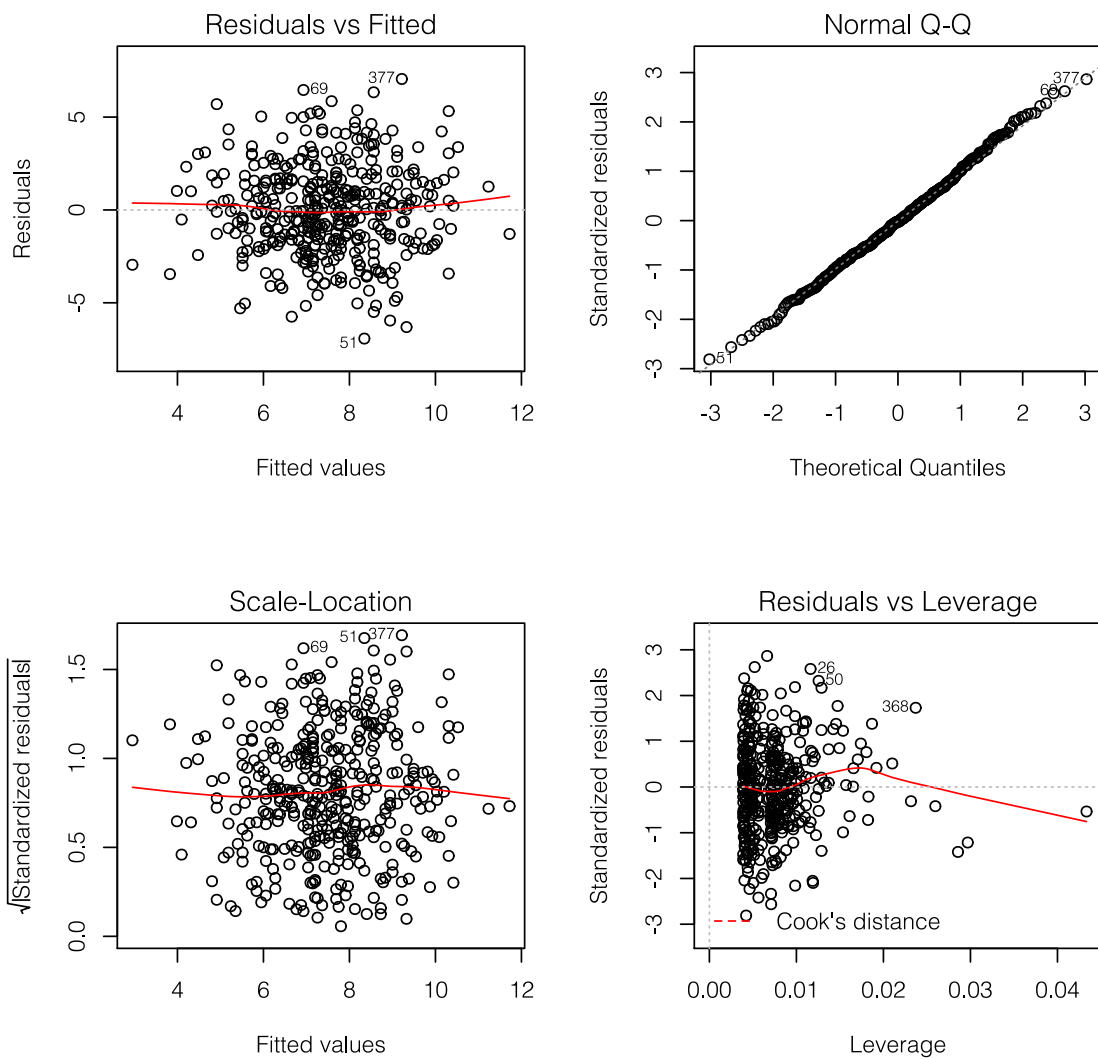
Part G

```
In [5]: confint(Carseats.mm.bs)
```

Out[5]:		2.5 %	97.5 %
	(Intercept)	11.79032	14.27127
	Price	-0.06475984	-0.04419543
	USYes	0.6915196	1.7077663

Part H

```
In [6]: par(mfrow = c(2,2))
plot(Carseats.mm.bs)
```



In terms of leverage, one of the points does seem to have a pull. However, there doesn't seem to be any noticeable outliers from the Residuals vs Fitted.

Exercise 4.7 #2

$$p_k(x) = \frac{\pi_k * \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{1}{2\sigma^2} (x - u_k)^2 \right\}}{\sum_{i=1}^K \pi_i * \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{1}{2\sigma^2} (x - u_i)^2 \right\}}$$

$$\delta_k(x) = x * \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k)$$

$\frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{1}{2\sigma^2} (x - u_i)^2 \right\}$ is constant for all k values, therefore, it does not factor into maximizing the entire function.

$$\log \left(\pi_k * \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{1}{2\sigma^2} (x - u_k)^2 \right\} \right) = \log(\pi_k) + \log \left(\frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{1}{2\sigma^2} (x - u_k)^2 \right\} \right) = \log(\pi_k) + \log\left(\frac{1}{\sqrt{2\pi}\sigma}\right) * \left(-\frac{1}{2\sigma^2} (x - u_k)^2\right)$$

$\log\left(\frac{1}{\sqrt{2\pi}\sigma}\right)$ is a constant multiple for all numerators that does not change between different k values; therefore, also should not factor into the maximization of the function. $\log(\pi_k) + -\frac{1}{2\sigma^2} (x^2 - 2xu_k + \mu_k^2)$

x^2 doesn't change based on different k-values, and therefore shouldn't factor into the maximizing function.

$$\log(\pi_k) + -\frac{1}{2\sigma^2} (-2xu_k + \mu_k^2) = \log(\pi_k) - \frac{1}{2\sigma^2} * -2xu_k - \frac{1}{2\sigma^2} * \mu_k^2$$

This is the same as:

$$\delta_k(x) = x * \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k)$$

Exercise 4.7 #5

Part A

For the training sets, QDA would perform better because the method is more flexible and could overcome issues of complexity within the training. For the testing sets, we would expect LDA to perform better than QDA if the Bayes Decision boundary is Linear because QDA might overfit on the training data set.

Part B

For the training and testing sets, we would expect LDA to perform better than QDA if the Bayes Decision boundary is non-Linear. This is because the non-Linear Bayes Decision boundary indicates great complexity within the data requiring the use of a more flexible method.

Part C

Given an extremely small n, it's possible that the QDA would be a better alternative to LDA because the Linearity of the Bayes Decision boundary could be questionable. However, this cannot conclude a general case.

Part D

False, QDA could overfit on the training data set which would affect the error rate on the test data set.

Exercise 4.7 #6

Part A

$$Y = -6 + 0.05 * Hours + GPA$$
$$p(X) = \frac{e^{-6+0.05*Hours+GPA}}{1 + e^{-6+0.05*Hours+GPA}}$$

```
In [7]: exp(-6 + 0.05*40 + 3.5)/(1+exp(-6 + 0.05*40 + 3.5))
```

```
Out[7]: 0.377540668798145
```

Part B

$$\frac{e^{-6+0.05*Hours+3.5}}{1 + e^{-6+0.05*Hours+3.5}} > 0.5 \implies$$
$$e^{-6+0.05*Hours+3.5} > 0.5 * (1 + e^{-6+0.05*Hours+3.5}) \implies$$
$$e^{-6+0.05*Hours+3.5} * 0.5 > 0.5 \implies$$
$$e^{-6+0.05*Hours+3.5} > 1 \implies$$
$$-6 + 0.05 * Hours + 3.5 > \log(1) \implies$$
$$0.05 * Hours > 0 + 2.5 \implies Hours > 2.5/0.05 = 50$$

Exercise 4.7 #11

Part A

```
In [8]: library(ISLR)
data(Auto)
Auto$mpg01 = rep(0, times = nrow(Auto))
Auto$mpg01[Auto$mpg>median(Auto$mpg)] = 1
```

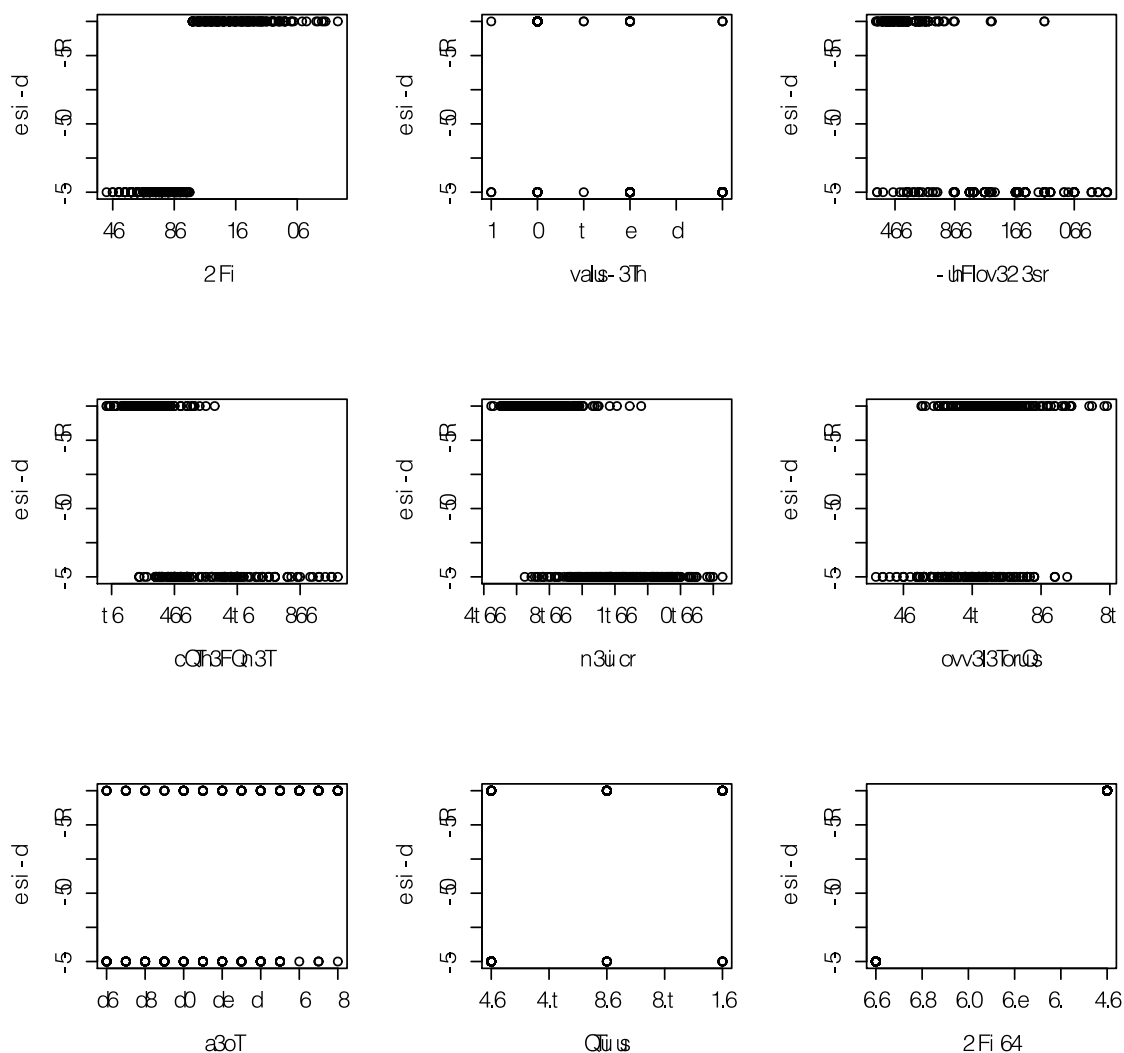
```
In [9]: cor(Auto[, -9])[, ncol(Auto[, -9])]
```

```
Out[9]:
```

mpg	0.836939229153618
cylinders	-0.75919388654222
displacement	-0.753476593523588
horsepower	-0.667052581047159
weight	-0.75775657142698
acceleration	0.346821530080934
year	0.429904226574927
origin	0.513698448318672
mpg01	1

Part B

```
In [10]: par(mfrow = c(3,3))
for(i in names(Auto[,-9])){
  plot(Auto[,i], Auto[, "mpg01"], xlab = i, ylab = "mpg01")
}
```



I couldn't really tell that much from the scatterplots, so I looked the correlations. From that mpg01 has best relation between cylinders, displacement, horsepower, and weight.

Part C

```
In [27]: training = sample(x = nrow(Auto), size = 0.5*nrow(Auto), replace = F)
training.set = Auto[training,]
test.set = Auto[-training,-10]
```

Part D

```
In [30]: library(MASS)
lda.auto = lda(mpg01~cylinders+displacement+horsepower+weight, data = training.set[, -9])
```

```
In [31]: lda.predict = predict(lda.auto, test.set)
mean(lda.predict$class != Auto[-training,10])
```

```
Out[31]: 0.112244897959184
```

Part E

```
In [32]: library(MASS)
qda.auto = qda(mpg01~cylinders+displacement+horsepower+weight, data = training.set[, -9])
```

```
In [33]: qda.predict = predict(qda.auto, test.set)
mean(qda.predict$class != Auto[-training,10])
```

```
Out[33]: 0.0969387755102041
```

Part F

```
In [35]: logistic.regrssion.auto = glm(mpg01~cylinders+displacement+horsepower+weight,
                                         data = training.set[, -9], family = binomial)
```

```
In [40]: logistic.regrssion.predict = predict(logistic.regrssion.auto, test.set, type="response")
logistic.regrssion.pred = rep(0, times = length(logistic.regrssion.predict))
logistic.regrssion.pred[logistic.regrssion.predict>0.5] = 1
mean(logistic.regrssion.pred != Auto[-training,10])
```

```
Out[40]: 0.112244897959184
```

Part G

```

In [56]: ks = 1:20
library(FNN)
k.results = data.frame(ks, error = ks)
for(i in ks){
  knn.predict = knn(train = training.set[,c("cylinders","displacement","horsepower","weight")],
    test = test.set[,c("cylinders","displacement","horsepower","weight")], cl = training.set$mpg01, k = i)
  k.results[k.results$ks == i, "error"] = mean(knn.predict != Auto[-training,10])
}
k.results

```

Out[56]:

	ks	error
1	1	0.1071429
2	2	0.1122449
3	3	0.1173469
4	4	0.122449
5	5	0.1173469
6	6	0.1122449
7	7	0.1122449
8	8	0.122449
9	9	0.1122449
10	10	0.1122449
11	11	0.1122449
12	12	0.1122449
13	13	0.1173469
14	14	0.122449
15	15	0.127551
16	16	0.127551
17	17	0.1326531
18	18	0.1326531
19	19	0.1326531
20	20	0.127551

The k values with the lowest error is k = 1.