

Analytics 512 Homework 4

Arif Ali

03/10/2015

Exercise 5.4 #3

Part A

K fold validation is implemented by first randomly dividing the sets of observations into k groups that should be about equal sizes. The first group (or fold) is treated as the validation set and the method is fit on the remaining k-1 folds. Then the mean-squared error (MSE) is computed for the fold, this method is then repeated on the remaining k-1 folds.

The CV estimate is computed by averaging the various MSEs

$$CV_k = \frac{1}{k} \sum_{i=1}^k MSE_i$$

Part B

i

The advantages of k fold over the validation set approach are:

- the validation estimate of the test error rate can be highly variable

- There's only one subset used to test and train, which makes it easy to implement

The disadvantages of k fold over the validation set approach are:

- The validation method is significantly easier to implement since we are splitting only two data sets

- The validation method has significantly less bias compared to the k-fold method.

ii

The disadvantages of k fold over LOOCV approach are:

The LOOCV is significantly more comprehensive compared to the k-fold method
has a low bias

The advantages of k fold over LOOCV approach are:

While LOOCV has a low bias, the test error can be highly variable as well

LOOCV requires going through every point, so it requires more time to implement than k-fold

Exercise 5.4 #5

```
In [35]: library(ISLR)
         attach(Default)
```

The following objects are masked from Default (pos = 3):

balance, default, income, student

The following objects are masked from Default (pos = 6):

balance, default, income, student

Part A

```
In [36]: glm.default = glm(default~balance+income, data = Default, family = binomial)
```

Part B

i

```
In [37]: train = sample(nrow(Default), .5*nrow(Default), replace = F)
         training = Default[train,]
         validation = Default[-train,]
```

ii

```
In [38]: glm.default = glm(default~., data = training, family = binomial)
```

iii

```
In [39]: pred = predict(glm.default, validation, type = "response")
pred.default = rep(x = "No", times = length(pred))
pred.default[pred>0.5] = "Yes"
```

iv

```
In [40]: mean(pred.default != validation$default)
```

```
Out[40]: 0.0292
```

I attempted CV by applying $K = 10$ instead of the validation method as specified by the homework.

```
In [41]: glm.default = glm(default~income+balance, data = Default, family = binomial)
cost <- function(r, pi = 0){
  mean(abs(r-pi) > 0.5)
}
library(boot)
cv.default = cv.glm(Default, glm.default, K = 10)
cv.default$delta
```

```
Out[41]: 0.0214718792280596 0.0214682053025623
```

Part C

```

In [42]: glm.default = glm(default~income+balance, data = Default, family = binomial)
cost <- function(r, pi = 0){
  mean(abs(r-pi) > 0.5)
}
library(boot)
cv.default = cv.glm(Default, glm.default, cost, K = 10)
cv.default$delta

glm.default = glm(default~income+balance, data = Default, family = binomial)
cost <- function(r, pi = 0){
  mean(abs(r-pi) > 0.5)
}
library(boot)
cv.default = cv.glm(Default, glm.default, cost, K = 10)
cv.default$delta

glm.default = glm(default~income+balance, data = Default, family = binomial)
cost <- function(r, pi = 0){
  mean(abs(r-pi) > 0.5)
}
library(boot)
cv.default = cv.glm(Default, glm.default, cost, K = 10)
cv.default$delta

```

```
Out[42]: 0.0263 0.02634
```

```
Out[42]: 0.0264 0.02642
```

```
Out[42]: 0.0265 0.02645
```

Based on the three runs of the, it seems that the CV error holds between 0.026 to 0.0265.

Part D

```

In [43]: glm.default = glm(default~income+balance+student, data = Default, family = binomial)
cost <- function(r, pi = 0){
  mean(abs(r-pi) > 0.5)
}
library(boot)
cv.default = cv.glm(Default, glm.default, cost, K = 10)
cv.default$delta

```

```
Out[43]: 0.0268 0.02694
```

The difference in error seems to be minimal, but lower without the student dummy variable. However, I would argue that the differences between the CV errors are small enough that it wouldn't matter. As to the specific question, adding the student dummy variable does not lead to a reduction in the test errors.

Exercise 6.8 #1

Part A

In the case of Best Subset selection, "This task must be performed with care, because the RSS of these $p + 1$ models decreases monotonically, and the R^2 increases monotonically, as the number of features included in the models increases." Thus Best Subset selection has the smallest training RSS.

Part B

Best Subset Selection would most likely have the smallest test RSS, but it's possible for forward stepwise or backwards stepwise to have a smaller test RSS, so the model with the smallest test RSS cannot be determined.

Part C

i

True, forward selection adds predictors to a model until it goes through each predictor, adding them if they meet the specified criteria.

ii

True, Backward selection starts by looking at all $k + 1$ predictors, removing ones that don't fit a criteria.

iii

False, the methods could lead to different models thus one cannot be a subset of the other.

iv

False, the methods could lead to different models thus one cannot be a subset of the other.

v

False, since best subset selection looks at all possible combinations, there are multiple models with $k+1$ predictors, so it doesn't have the same values as the lower subsets.

Exercise 6.8 #8

Part A

```
In [44]: X = rnorm(100)
         e = rnorm(100)
```

Part B

```
In [45]: Y = 1 +2*X+3*X^2+4*X^3+e
```




Part C

```
In [46]: XY = data.frame(Y,X)
names(XY)
library(leaps)
best_selection = regsubsets(Y~poly(X, 10), data = XY,nvmax=10)
```

```
Out[46]:      'Y'  'X'
```

```
In [47]: best_selection_summary = summary(best_selection)
```

```
In [48]: par(mfrow = c(3,1))
plot(best_selection_summary$cp, type = "b", xlab = "degree of polynomial" , ylab = "Cp")
plot(best_selection_summary$bic,type = "b", xlab = "degree of polynomial" , ylab = "bic")
plot(best_selection_summary$rsq, type = "b", xlab = "degree of polynomial" , ylab = "rsq")
```



```
In [49]: print("Based on Cp, the best model is:")
coef(best_selection, which.min(best_selection_summary$cp))
print("Based on BIC, the best model is:")
coef(best_selection, which.min(best_selection_summary$bic))
print("Based on R squared, the best model is:")
coef(best_selection, which.max(best_selection_summary$rsq))
```

```
[1] "Based on Cp, the best model is:"
```

```
Out[49]:      (Intercept)  4.56470622794346
      poly(X, 10)1 149.309802114251
      poly(X, 10)2 38.8087847420602
      poly(X, 10)3 84.9219405421287
      poly(X, 10)4 -2.28234848572183
```

```
[1] "Based on BIC, the best model is:"
```




```
Out[49]:      (Intercept)  4.56470622794346
      poly(X, 10)1 149.309802114251
      poly(X, 10)2 38.8087847420602
      poly(X, 10)3 84.9219405421287
      poly(X, 10)4 -2.28234848572183
```

```
[1] "Based on R squared, the best model is:"
```

```
Out[49]:      (Intercept)  4.56470622794346
      poly(X, 10)1 149.309802114251
      poly(X, 10)2 38.8087847420602
      poly(X, 10)3 84.9219405421287
      poly(X, 10)4 -2.28234848572183
      poly(X, 10)5 0.457699187806449
      poly(X, 10)6 1.28566571981932
      poly(X, 10)7 -0.107867813523064
      poly(X, 10)8 0.863027733541585
      poly(X, 10)9 0.113331849816744
      poly(X, 10)10 0.149897834460875
```

Part D


```
In [50]: forward_selection = regsubsets(Y~poly(X, 10), data = XY, method = "forward")
forward_selection_summary = summary(forward_selection)
par(mfrow = c(3,1))
plot(forward_selection_summary$cp,type = "b", xlab = "degree of polynomial" ,
ylab = "Cp")
plot(forward_selection_summary$bic,type = "b", xlab = "degree of polynomial" ,
ylab = "bic")
plot(forward_selection_summary$rsq,type = "b", xlab = "degree of polynomial" ,
ylab = "rsq")
```



```
In [51]: print("Based on Cp, the best model is:")
coef(forward_selection,which.min(forward_selection_summary$cp))
print("Based on BIC, the best model is:")
coef(forward_selection,which.min(forward_selection_summary$bic))
print("Based on R squared, the best model is:")
coef(forward_selection,which.max(forward_selection_summary$rsq))
```

```
[1] "Based on Cp, the best model is:"
```

```
Out[51]:      (Intercept)  4.56470622794346
      poly(X, 10)1 149.309802114251
      poly(X, 10)2 38.8087847420602
      poly(X, 10)3 84.9219405421287
      poly(X, 10)4 -2.28234848572183
```




```
[1] "Based on BIC, the best model is:"
```

```
Out[51]:      (Intercept)  4.56470622794346
      poly(X, 10)1 149.309802114251
      poly(X, 10)2 38.8087847420602
      poly(X, 10)3 84.9219405421287
      poly(X, 10)4 -2.28234848572183
```

```
[1] "Based on R squared, the best model is:"
```

```
Out[51]:      (Intercept)  4.56470622794346
      poly(X, 10)1 149.309802114251
      poly(X, 10)2 38.8087847420602
      poly(X, 10)3 84.9219405421287
      poly(X, 10)4 -2.28234848572183
      poly(X, 10)5 0.457699187806449
      poly(X, 10)6 1.28566571981932
      poly(X, 10)8 0.863027733541585
      poly(X, 10)10 0.149897834460875
```

```
In [52]: backward_selection = regsubsets(Y~poly(X, 10), data = XY, method = "backward")
backward_selection_summary = summary(backward_selection)
par(mfrow = c(3,1))
plot(backward_selection_summary$cp,type = "b", xlab = "degree of polynomial" ,
ylab = "Cp")
plot(backward_selection_summary$bic,type = "b", xlab = "degree of polynomial"
, ylab = "bic")
plot(backward_selection_summary$rsq,type = "b", xlab = "degree of polynomial"
, ylab = "rsq")
```



```
In [53]: print("Based on Cp, the best model is:")
coef(backward_selection,which.min(backward_selection_summary$cp))
print("Based on BIC, the best model is:")

coef(backward_selection,which.min(backward_selection_summary$bic))
print("Based on R squared, the best model is:")

coef(backward_selection,which.max(backward_selection_summary$rsq))
```

```
[1] "Based on Cp, the best model is:"
```

```
Out[53]:      (Intercept)  4.56470622794346
      poly(X, 10)1 149.309802114251
      poly(X, 10)2 38.8087847420602
      poly(X, 10)3 84.9219405421287
      poly(X, 10)4 -2.28234848572183
```

```
[1] "Based on BIC, the best model is:"
```

```
Out[53]:      (Intercept)  4.56470622794346
      poly(X, 10)1 149.309802114251
      poly(X, 10)2 38.8087847420602
      poly(X, 10)3 84.9219405421287
      poly(X, 10)4 -2.28234848572183
```

```
[1] "Based on R squared, the best model is:"
```

```
Out[53]:      (Intercept)  4.56470622794346
      poly(X, 10)1 149.309802114251
      poly(X, 10)2 38.8087847420602
      poly(X, 10)3 84.9219405421287
      poly(X, 10)4 -2.28234848572183
      poly(X, 10)5 0.457699187806449
      poly(X, 10)6 1.28566571981932
      poly(X, 10)8 0.863027733541585
      poly(X, 10)10 0.149897834460875
```

All the models based on R^2 , Cp, and BIC seem to conclude the same model types. Both Cp and BIC pick models with a polynomial degree of 4 whereas R^2 picks a 10th degree polynomial model.